



Feature Store & Batch Inference: AI-Powered Apple Leaf Specialist

- **Name:** Aktham Almomani
- **Group:** 7

Introduction

This notebook builds a small, end-to-end feature store workflow for the apple leaf project. It creates a feature group for image metadata, ingests a demo table, validates online and offline stores, runs Athena queries, and demonstrates batch inference with a PyTorch model using JSONL input.

Here're the main steps in this notebooks:

- Stand up a Feature Group for image-level features.
- Ingest and validate records in online and offline stores.
- Query features via Athena for analysis.
- Run a SageMaker Batch Transform job with an existing model artifact.

Setup and Parameters

First, let's setup using SageMaker SDK to obtain session, role, region, and default S3 bucket. Bucket paths are used for the offline store and Athena outputs

```
In [ ]: import boto3, sagemaker, pandas as pd, time, uuid
from sagemaker.feature_store.feature_group import FeatureGroup
from sagemaker.feature_store.feature_definition import FeatureDefinition, FeatureGroupDefinition

sess = sagemaker.Session()
role = sagemaker.get_execution_role()
region = sess.boto_region_name
bucket = sess.default_bucket()

fg_name = "apple_image_meta_fg"
offline_s3_uri = f"s3://{bucket}/feature-store/{fg_name}/offline"

# here let's define the schema:
feature_definitions = [
    FeatureDefinition(feature_name="image_id", feature_type=FeatureType.STRING),
    FeatureDefinition(feature_name="s3_uri", feature_type=FeatureType.STRING),
    FeatureDefinition(feature_name="split", feature_type=FeatureType.STRING),
    FeatureDefinition(feature_name="label", feature_type=FeatureType.STRING),
    FeatureDefinition(feature_name="brightness_bin", feature_type=FeatureType.STRING)
```

```

FeatureDefinition(feature_name="avg_r", feature_type=FeatureType
FeatureDefinition(feature_name="avg_g", feature_type=FeatureType
FeatureDefinition(feature_name="avg_b", feature_type=FeatureType
FeatureDefinition(feature_name="event_time", feature_type=FeatureType
]

```

sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml

sagemaker.config INFO - Not applying SDK defaults from location: /home/sagemaker-user/.config/sagemaker/config.yaml

Define Feature Group

Here let's create `apple_image_meta_fg` with schema fields such as `image_id`, `s3_uri`, `split`, `label`, `brightness_bin`, `avg_r`, `avg_g`, `avg_b`, and `event_time`. Configure `OnlineStore` enabled and `OfflineStore` to S3. Handle SDK version differences for `create()` signatures.

- **brightness_bin:** A categorical bucket that encodes overall image brightness (luminance).
- **avg_r:** Mean intensity of the `red` channel across all pixels in the image.
- **avg_g:** Mean intensity of the `green` channel across all pixels.
- **avg_b:** Mean intensity of the `blue` channel across all pixels.

```

In [ ]: fg = FeatureGroup(name=fg_name, sagemaker_session=sess)

try:
    fg.delete()
except Exception:
    pass

try:
    fg.feature_definitions = feature_definitions
except Exception:
    pass

try:
    fg.create(
        feature_definitions,
        record_identifier_name="image_id",
        event_time_feature_name="event_time",
        role_arn=role,
        online_store_config={"EnableOnlineStore": True},
        offline_store_config={"S3StorageConfig": {"S3Uri": offline_s3_uri}},
    )
except TypeError:

```

```

fg.create(
    s3_uri=offline_s3_uri,
    record_identifier_name="image_id",
    event_time_feature_name="event_time",
    role_arn=role,
    enable_online_store=True,
)

print("Feature group created:", fg.describe()["FeatureGroupName"])

```

Feature group created: apple_image_meta_fg

Ingest Demo Records

Here let's build a tiny DataFrame with two sample images and matching columns, then call `fg.ingest(..., wait=True)` to load records.

```

In [ ]: from datetime import datetime, timezone

now = datetime.now(timezone.utc).isoformat()
demo = [
    {
        "image_id": str(uuid.uuid4()),
        "s3_uri": f"s3://{bucket}/apple/sample/healthy1.jpg",
        "split": "test",
        "label": "healthy",
        "brightness_bin": "bright",
        "avg_r": 0.52,
        "avg_g": 0.61,
        "avg_b": 0.49,
        "event_time": now,
    },
    {
        "image_id": str(uuid.uuid4()),
        "s3_uri": f"s3://{bucket}/apple/sample/rust1.jpg",
        "split": "test",
        "label": "rust",
        "brightness_bin": "dim",
        "avg_r": 0.42,
        "avg_g": 0.50,
        "avg_b": 0.55,
        "event_time": now,
    },
]
df = pd.DataFrame(demo)
cols = ["image_id", "s3_uri", "split", "label", "brightness_bin", "avg_r", "avg_g", "avg_b", "event_time"]
df = df[cols]
df

```

Out[]:	image_id	s3_uri	split	label	brightness_bi
0	b544-43a8-abfa-63247d39ca22	91c1fb8d- east-1-533266958221/ apple/sa...	test	healthy	brigl
1	1a37c753-affc-46ff- b902-3fc20832620d	s3://sagemaker-us- east-1-533266958221/ apple/sa...	test	rust	di

Feature Group Status Check

here let's poll `fg.describe()["FeatureGroupStatus"]` to ensure the group is Created before downstream operations.

```
In [ ]: def wait_fg_active(fg, poll=5):
        while True:
            status = fg.describe()["FeatureGroupStatus"]
            print("FeatureGroupStatus:", status)
            if status in ("Created", "CreateFailed"):
                break
            time.sleep(poll)

        wait_fg_active(fg)
```

FeatureGroupStatus: Created

Query Offline Store with Athena

Use Feature Store query helpers to create and run SQL against the Glue Data Catalog database and table backing the offline store. Fetch results as a DataFrame for quick validation and simple aggregations.

```
In [ ]: # Here's let's query the offline store via Athena:
        from sagemaker.feature_store.feature_store import FeatureStore

        fs = FeatureStore(sagemaker_session=sess)
        query = fs.create_query_with_feature_group(fg_name, table_name=None)
        sql = f"SELECT image_id, label, brightness_bin, avg_r, avg_g, avg_b FROM {quer
        athena = query.run(sql_query=sql, output_location=f"s3://{bucket}/athena-resul
        athena.as_dataframe().head()
```

```
In [ ]: ingest_res = fg.ingest(data_frame=df, max_workers=4, wait=True)
        print("Ingest complete.")
```

Ingest complete.

Online Store Read

Verify near-real-time availability with `sagemaker-featurestore-runtime` `get_record` for a sample `image_id`, confirming keys and values returned from the online store.

```
In [ ]: fs_rt = boto3.client("sagemaker-featurestore-runtime")

one_id = df.iloc[0]["image_id"]
resp = fs_rt.get_record(
    FeatureGroupName=fg.name,
    RecordIdentifierValueAsString=str(one_id)
)
print("Online GetRecord keys:", [f["FeatureName"] for f in resp.get("Record",
```

Online GetRecord keys: ['image_id', 's3_uri', 'split', 'label', 'brightness_bin', 'avg_r', 'avg_g', 'avg_b', 'event_time']

```
In [ ]: desc = fg.describe()
print("OfflineStoreConfig:", desc.get("OfflineStoreConfig"))
print("OfflineStoreStatus:", desc.get("OfflineStoreStatus"))
print("DataCatalogConfig:", desc.get("OfflineStoreConfig", {}).get("DataCatalogConfig"))
print("S3 uri:", desc.get("OfflineStoreConfig", {}).get("S3StorageConfig"))
```

```
OfflineStoreConfig: {'S3StorageConfig': {'S3Uri': 's3://sagemaker-us-east-1-533266958221/feature-store/apple_image_meta_fg/offline', 'ResolvedOutputS3Uri': 's3://sagemaker-us-east-1-533266958221/feature-store/apple_image_meta_fg/offline/533266958221/sagemaker/us-east-1/offline-store/apple_image_meta_fg-1760761969/data'}, 'DisableGlueTableCreation': False, 'DataCatalogConfig': {'TableName': 'apple_image_meta_fg_1760761969', 'Catalog': 'AwsDataCatalog', 'Database': 'sagemaker_featurestore'}}
OfflineStoreStatus: {'Status': 'Active'}
DataCatalogConfig: {'TableName': 'apple_image_meta_fg_1760761969', 'Catalog': 'AwsDataCatalog', 'Database': 'sagemaker_featurestore'}
S3 uri: {'S3Uri': 's3://sagemaker-us-east-1-533266958221/feature-store/apple_image_meta_fg/offline', 'ResolvedOutputS3Uri': 's3://sagemaker-us-east-1-533266958221/feature-store/apple_image_meta_fg/offline/533266958221/sagemaker/us-east-1/offline-store/apple_image_meta_fg-1760761969/data'}
```

```
In [ ]: from sagemaker.feature_store.feature_group import AthenaQuery
import json

# Pull the real DB/table from describe()
desc = fg.describe()
db = desc["OfflineStoreConfig"]["DataCatalogConfig"]["Database"]
tbl = desc["OfflineStoreConfig"]["DataCatalogConfig"]["TableName"]
print("Using:", db, ".", tbl)

# Create the helper and point it explicitly at the DB & table
query = fg.athena_query()
query.database = db
query.table_name = tbl
```

```

# Pick an S3 location for Athena results:
output_loc = f"s3://{bucket}/athena/{tbl}/"

# Run a simple aggregation
sql = f'SELECT label, COUNT(*) AS n FROM "{db}"."{tbl}" GROUP BY label'
query.run(query_string=sql, output_location=output_loc)
query.wait()

# Fetch results as a dataframe
athena_df = query.as_dataframe()
athena_df

```

Using: sagemaker_featurestore . apple_image_meta_fg_1760761969

```

Out[ ]:
   label  n
0    rust  1
1  healthy  1

```

Batch Inference (Transform Job)

Now let's prepare a JSONL file of base64-encoded images from `val.zip`, upload to S3, and launch a `PyTorchModel.transformer` job that uses your packaged `model.tar.gz` and `inference.py`. Wait for completion and capture S3 output location.

```

In [ ]: # Batch Transform using the trained estimator:

```

```

import base64, os
from pathlib import Path
from zipfile import ZipFile
from datetime import datetime
from sagemaker.pytorch import PyTorchModel

sess = sagemaker.Session()
role = sagemaker.get_execution_role()
bucket = sess.default_bucket()
s3 = boto3.client("s3")

# let's start by getting val.zip:
zip_key = "apple/val.zip"
local_zip = "/tmp/val.zip"
s3.download_file(bucket, zip_key, local_zip)

```

```

In [ ]: # Extract a handful of images:
extract_dir = Path("/tmp/val_images")
extract_dir.mkdir(exist_ok=True, parents=True)
with ZipFile(local_zip) as zf:
    zf.extractall(extract_dir)

```

```

# Collect some image files:
img_paths = []
for ext in (".jpg", ".jpeg", ".png"):
    img_paths += list(extract_dir.rglob(f"*{ext}"))
img_paths = img_paths[:20]

assert img_paths, "No images found after unzipping! Check the archive structure"

```

```

In [ ]: # Build JSONL:
local_jsonl = Path("/tmp/batch_input.jsonl")
with local_jsonl.open("w") as f:
    for p in img_paths:
        with open(p, "rb") as im:
            b64 = base64.b64encode(im.read()).decode("utf-8")
            f.write(json.dumps({"b64": b64}) + "\n")

```

```

In [ ]: # Upload JSONL to S3:
s3_in = sess.upload_data(str(local_jsonl), bucket=bucket, key_prefix="batch/input")
print("Uploaded to:", s3_in)

```

Uploaded to: s3://sagemaker-us-east-1-533266958221/batch/input/batch_input.jsonl

```

In [ ]: # Finally let's run Batch Transform using your PyTorch inference code:

# Our trained model.tar.gz:
model_artifact = "s3://sagemaker-us-east-1-533266958221/apple/outputs/pytorch-

pyt_model = PyTorchModel(
    model_data=model_artifact,
    role=role,
    framework_version="2.3",
    py_version="py311",
    entry_point="inference.py",
    source_dir="src",
    sagemaker_session=sess,
)

run_id = datetime.utcnow().strftime("%Y%m%d-%H%M%S")
s3_out = f"s3://{bucket}/batch/output/{run_id}/"

transformer = pyt_model.transformer(
    instance_count=1,
    instance_type="ml.m5.large",
    output_path=s3_out,
    strategy="SingleRecord",
)

transformer.transform(
    data=s3_in,
    content_type="application/json",
    split_type="Line",

```

```
)  
transformer.wait()  
  
print("Batch input :", s3_in)  
print("Batch output:", s3_out)
```

```
/tmp/ipykernel_3627/440194388.py:15: DeprecationWarning: datetime.datetime.utcnow()  
ow() is deprecated and scheduled for removal in a future version. Use timezone-  
aware objects to represent datetimes in UTC: datetime.datetime.now(datetime.UT  
C).
```

```
run_id = datetime.utcnow().strftime("%Y%m%d-%H%M%S")  
INFO:sagemaker:Repacking model artifact (s3://sagemaker-us-east-1-533266958221/  
apple/outputs/pytorch-training-2025-10-16-04-42-30-265/output/model.tar.gz), sc  
ript artifact (src), and dependencies ([]) into single tar.gz file located at s  
3://sagemaker-us-east-1-533266958221/pytorch-inference-2025-10-18-06-06-31-168/  
model.tar.gz. This may take some time depending on model size...  
INFO:sagemaker:Creating model with name: pytorch-inference-2025-10-18-06-06-3  
5-301  
INFO:sagemaker:Creating transform job with name: pytorch-inference-2025-10-18-0  
6-06-36-031
```

```
.....['torchserve', '--start', '--model-store', '/.sagemaker/ts/models', '--ts-config', '/etc/sagemaker-ts.properties', '--log-config', '/opt/conda/lib/python3.11/site-packages/sagemaker_pytorch_serving_container/etc/log4j2.xml', '--models', 'model=/opt/ml/model']
Warning: TorchServe is using non-default JVM parameters: -XX:-UseContainerSupport
WARNING: sun.reflect.Reflection.getCallerClass is not supported. This will impact performance.
2025-10-18T06:11:50,317 [WARN ] main org.pytorch.serve.util.ConfigManager - Your torchserve instance can access any URL to load models. When deploying to production, make sure to limit the set of allowed_urls in config.properties
2025-10-18T06:11:50,321 [INFO ] main org.pytorch.serve.servingsdk.impl.PluginsManager - Initializing plugins manager...
2025-10-18T06:11:50,402 [INFO ] main org.pytorch.serve.metrics.configuration.MetricConfiguration - Successfully loaded metrics configuration from /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml
2025-10-18T06:11:50,499 [INFO ] main org.pytorch.serve.ModelServer - Torchserve version: 0.11.0
TS Home: /opt/conda/lib/python3.11/site-packages
Current directory: /
Temp directory: /tmp
Metrics config path: /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml
Number of GPUs: 0
Number of CPUs: 2
Max heap size: 1918 M
Python executable: /opt/conda/bin/python3.11
Config file: /etc/sagemaker-ts.properties
Inference address: http://0.0.0.0:8080
Management address: http://0.0.0.0:8080
Metrics address: http://127.0.0.1:8082
Model Store: /.sagemaker/ts/models
Initial Models: model=/opt/ml/model
Log dir: /logs
Metrics dir: /logs
Netty threads: 0
Netty client threads: 0
Default workers per model: 2
Blacklist Regex: N/A
Maximum Response Size: 6553500
Maximum Request Size: 6553500
Limit Maximum Image Pixels: true
Prefer direct buffer: false
Allowed Urls: [file://.*|http(s)?://.*]
Custom python dependency for model allowed: false
Enable metrics API: true
Metrics mode: LOG
Disable system metrics: false
Workflow Store: /.sagemaker/ts/models
CPP log config: N/A
Model config: N/A
System metrics command: default
2025-10-18T06:11:50,509 [INFO ] main org.pytorch.serve.servingsdk.impl.PluginsManager - Loading snapshot serializer plugin...
```

```
2025-10-18T06:11:50,534 [INFO ] main org.pytorch.serve.ModelServer - Loading in
ital models: /opt/ml/model
2025-10-18T06:11:50,546 [INFO ] main org.pytorch.serve.archive.model.ModelArchi
ve - createTempDir /tmp/models/99953fda636149fdbca0a2785a7884fc
2025-10-18T06:11:50,546 [INFO ] main org.pytorch.serve.archive.model.ModelArchi
ve - createSymbolicDir /tmp/models/99953fda636149fdbca0a2785a7884fc/model
2025-10-18T06:11:50,548 [WARN ] main org.pytorch.serve.archive.model.ModelArchi
ve - Model archive version is not defined. Please upgrade to torch-model-archiv
er 0.2.0 or higher
2025-10-18T06:11:50,548 [WARN ] main org.pytorch.serve.archive.model.ModelArchi
ve - Model archive createdOn is not defined. Please upgrade to torch-model-arch
iver 0.2.0 or higher
2025-10-18T06:11:50,551 [INFO ] main org.pytorch.serve.wlm.ModelManager - Model
model loaded.
2025-10-18T06:11:50,570 [INFO ] main org.pytorch.serve.ModelServer - Initialize
Inference server with: EpollServerSocketChannel.
2025-10-18T06:11:50,738 [INFO ] main org.pytorch.serve.ModelServer - Inference
API bind to: http://0.0.0.0:8080
2025-10-18T06:11:50,739 [INFO ] main org.pytorch.serve.ModelServer - Initialize
Metrics server with: EpollServerSocketChannel.
2025-10-18T06:11:50,753 [INFO ] main org.pytorch.serve.ModelServer - Metrics AP
I bind to: http://127.0.0.1:8082
Model server started.
2025-10-18T06:11:51,484 [WARN ] pool-3-thread-1 org.pytorch.serve.metrics.Metri
cCollector - worker pid is not available yet.
2025-10-18T06:11:51,672 [INFO ] pool-3-thread-1 TS_METRICS - CPUUtilization.Per
cent:100.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,673 [INFO ] pool-3-thread-1 TS_METRICS - DiskAvailable.Giga
bytes:47.07512283325195|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,674 [INFO ] pool-3-thread-1 TS_METRICS - DiskUsage.Gigabyte
s:8.790008544921875|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,674 [INFO ] pool-3-thread-1 TS_METRICS - DiskUtilization.Pe
rcent:15.7|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,675 [INFO ] pool-3-thread-1 TS_METRICS - MemoryAvailable.Me
gabytes:6569.18359375|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,675 [INFO ] pool-3-thread-1 TS_METRICS - MemoryUsed.Megabyt
es:803.8046875|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,675 [INFO ] pool-3-thread-1 TS_METRICS - MemoryUtilizatio
n.Percent:14.3|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:53,294 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - s_name_part
0=/tmp/.ts.sock, s_name_part1=9000, pid=36
2025-10-18T06:11:53,295 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - Listening o
n port: /tmp/.ts.sock.9000
2025-10-18T06:11:53,306 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - Successfull
y loaded /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml.
2025-10-18T06:11:53,307 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - [PID]36
2025-10-18T06:11:53,307 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - Torch worke
r started.
2025-10-18T06:11:53,308 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - Python runt
ime: 3.11.9
2025-10-18T06:11:53,314 [INFO ] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Connecting to: /tmp/.ts.sock.9000
2025-10-18T06:11:53,341 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - Connection
accepted: /tmp/.ts.sock.9000.
```

2025-10-18T06:11:53,364 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767913364
2025-10-18T06:11:53,545 [INFO] W-9000-model_1.0-stdout MODEL_LOG - model_name: model, batchSize: 1
2025-10-18T06:11:53,617 [INFO] W-9001-model_1.0-stdout MODEL_LOG - s_name_part0=/tmp/.ts.sock, s_name_part1=9001, pid=37
2025-10-18T06:11:53,620 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Listening on port: /tmp/.ts.sock.9001
2025-10-18T06:11:53,630 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Successfully loaded /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml.
2025-10-18T06:11:53,631 [INFO] W-9001-model_1.0-stdout MODEL_LOG - [PID]37
2025-10-18T06:11:53,632 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Torch worker started.
2025-10-18T06:11:53,633 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Python runtime: 3.11.9
2025-10-18T06:11:53,633 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Connecting to: /tmp/.ts.sock.9001
2025-10-18T06:11:53,641 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Connection accepted: /tmp/.ts.sock.9001.
2025-10-18T06:11:53,653 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767913653
2025-10-18T06:11:53,733 [INFO] W-9001-model_1.0-stdout MODEL_LOG - model_name: model, batchSize: 1
2025-10-18T06:11:55,863 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 2209
2025-10-18T06:11:55,864 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 2498
2025-10-18T06:11:55,864 [INFO] W-9001-model_1.0 TS_METRICS - WorkerLoadTime.Milliseconds:5299.0|#WorkerName:W-9001-model_1.0,Level:Host|#hostname:0da699790d1b,timestamp:1760767915
2025-10-18T06:11:55,865 [INFO] W-9000-model_1.0 TS_METRICS - WorkerLoadTime.Milliseconds:5305.0|#WorkerName:W-9000-model_1.0,Level:Host|#hostname:0da699790d1b,timestamp:1760767915
2025-10-18T06:11:55,866 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:17.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767915
2025-10-18T06:11:55,866 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:15.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767915
2025-10-18T06:12:02,600 [INFO] pool-2-thread-3 ACCESS_LOG - /169.254.255.130:36970 "GET /ping HTTP/1.1" 200 32
2025-10-18T06:12:02,601 [INFO] pool-2-thread-3 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767922
2025-10-18T06:12:02,692 [INFO] epollEventLoopGroup-3-2 ACCESS_LOG - /169.254.255.130:36978 "GET /execution-parameters HTTP/1.1" 404 4
2025-10-18T06:12:02,694 [INFO] epollEventLoopGroup-3-2 TS_METRICS - Requests4XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767922
2025-10-18T06:12:03,208 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,216 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923216
2025-10-18T06:12:03,219 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923
2025-10-18T06:12:03,368 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:148.52|#ModelName:

e:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,eb22ad0c-d5c7-416e-854c-8a047ae4898f, pattern=[METRICS]

2025-10-18T06:12:03,371 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:148.52|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:eb22ad0c-d5c7-416e-854c-8a047ae4898f,timestamp:1760767923

2025-10-18T06:12:03,373 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 169

2025-10-18T06:12:03,374 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,375 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:157701.166|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,375 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:159.077|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,378 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,378 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 156

2025-10-18T06:12:03,379 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,450 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,451 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923451

2025-10-18T06:12:03,457 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:03,565 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:112.06|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,00d978d5-99ef-4310-9fda-9399cce91dbb, pattern=[METRICS]

2025-10-18T06:12:03,566 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 117

2025-10-18T06:12:03,567 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:112.06|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:00d978d5-99ef-4310-9fda-9399cce91dbb,timestamp:1760767923

2025-10-18T06:12:03,568 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,569 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:115682.524|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,570 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:110.69|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,571 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,571 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 115

2025-10-18T06:12:03,572 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:7.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,586 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,588 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923588
2025-10-18T06:12:03,589 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923
2025-10-18T06:12:03,666 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:76.47|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,535aaf8f-12c2-4b25-aaf1-4903e128377d, pattern=[METRICS]
2025-10-18T06:12:03,667 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:76.47|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:535aaf8f-12c2-4b25-aaf1-4903e128377d,timestamp:1760767923
2025-10-18T06:12:03,667 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:03,668 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,669 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79881.416|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,670 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:118.211|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,671 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,671 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 79
2025-10-18T06:12:03,671 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:5.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,679 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,680 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923680
2025-10-18T06:12:03,681 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923
2025-10-18T06:12:03,759 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:77.19|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,139c1d02-7d60-448c-9b37-15386707603b, pattern=[METRICS]
2025-10-18T06:12:03,760 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:03,760 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,760 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:77.19|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:139c1d02-7d60-448c-9b37-15386707603b,timestamp:1760767923
2025-10-18T06:12:03,761 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79681.473|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,761 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:162.982|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,762 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,762 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread

read - Backend response time: 78
2025-10-18T06:12:03,762 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,775 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,776 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923776
2025-10-18T06:12:03,778 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923
2025-10-18T06:12:03,855 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:76.6|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,55b975c1-dc4b-46b1-92fc-af89488e5ce7, pattern=[METRICS]
2025-10-18T06:12:03,856 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:03,857 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,857 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:76.6|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:55b975c1-dc4b-46b1-92fc-af89488e5ce7,timestamp:1760767923
2025-10-18T06:12:03,858 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79770.92|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,858 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:104.25|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,859 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,860 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 79
2025-10-18T06:12:03,862 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:7.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,870 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,870 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923870
2025-10-18T06:12:03,873 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923
2025-10-18T06:12:03,955 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:82.16|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,4416b3f4-bd1b-4485-bc1d-448e7bf30711, pattern=[METRICS]
2025-10-18T06:12:03,955 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 86
2025-10-18T06:12:03,956 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,956 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:82.16|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:4416b3f4-bd1b-4485-bc1d-448e7bf30711,timestamp:1760767923
2025-10-18T06:12:03,956 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:85149.584|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,957 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:107.279|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,957 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,957 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 84

2025-10-18T06:12:03,958 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,967 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,968 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923968

2025-10-18T06:12:03,974 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:04,052 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:78.89|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,be1244e1-4fb2-4c32-8c4c-2e9803e5e3e9, pattern=[METRICS]

2025-10-18T06:12:04,053 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 87

2025-10-18T06:12:04,053 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,054 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:85517.614|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,054 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:97.436|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,055 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,055 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 84

2025-10-18T06:12:04,055 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,054 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:78.89|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:be1244e1-4fb2-4c32-8c4c-2e9803e5e3e9,timestamp:1760767924

2025-10-18T06:12:04,069 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,070 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924070

2025-10-18T06:12:04,073 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.59|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,6463fc01-0332-4cde-a248-a3d8a9620b7a, pattern=[METRICS]

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.59|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:6463fc01-0332-4cde-a248-a3d8a9620b7a,timestamp:1760767924

2025-10-18T06:12:04,150 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77731.709|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,150 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:103.801|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,150 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,151 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 78

2025-10-18T06:12:04,151 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:3.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,165 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,167 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924167

2025-10-18T06:12:04,169 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,245 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:76.2|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,db313afb-9d9e-4212-a11c-6b71df7d258b, pattern=[METRICS]

2025-10-18T06:12:04,245 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 80

2025-10-18T06:12:04,246 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,246 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:76.2|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:db313afb-9d9e-4212-a11c-6b71df7d258b,timestamp:1760767924

2025-10-18T06:12:04,247 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79525.992|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,247 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:220.372|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,248 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,248 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76

2025-10-18T06:12:04,248 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,260 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,261 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924261

2025-10-18T06:12:04,262 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,339 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.89|#modelName:m

odel,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,b5b61720-5375-4580-b0c8-2b0354434701, pattern=[METRICS]

2025-10-18T06:12:04,339 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79

2025-10-18T06:12:04,340 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,340 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.89|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:b5b61720-5375-4580-b0c8-2b0354434701,timestamp:1760767924

2025-10-18T06:12:04,341 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77997.837|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,342 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:221.097|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,342 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,342 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 77

2025-10-18T06:12:04,343 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:5.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,353 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,354 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924354

2025-10-18T06:12:04,355 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,437 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:81.17|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,876916ed-83a6-47bc-99eb-c38dcdf4ca2, pattern=[METRICS]

2025-10-18T06:12:04,437 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:81.17|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:876916ed-83a6-47bc-99eb-c38dcdf4ca2,timestamp:1760767924

2025-10-18T06:12:04,440 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 85

2025-10-18T06:12:04,440 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:84146.256|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:97.85|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 84

2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:3.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,454 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,456 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924456

2025-10-18T06:12:04,461 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,534 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.86|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,2c08a9c5-b823-4cf7-b907-546f89204feb, pattern=[METRICS]

2025-10-18T06:12:04,534 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.86|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:2c08a9c5-b823-4cf7-b907-546f89204feb,timestamp:1760767924

2025-10-18T06:12:04,535 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81

2025-10-18T06:12:04,536 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,536 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79458.604|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,536 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:113.16|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,537 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,537 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 78

2025-10-18T06:12:04,537 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,546 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,548 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924548

2025-10-18T06:12:04,550 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,625 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.4|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,99cf3977-ef89-4533-91d6-d8c24f50c4f3, pattern=[METRICS]

2025-10-18T06:12:04,626 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 80

2025-10-18T06:12:04,626 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.4|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:99cf3977-ef89-4533-91d6-d8c24f50c4f3,timestamp:1760767924

2025-10-18T06:12:04,627 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,627 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79242.782|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,628 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:110.513|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,628 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,629 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread

read - Backend response time: 76
2025-10-18T06:12:04,629 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,637 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,639 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924638
2025-10-18T06:12:04,640 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924
2025-10-18T06:12:04,714 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:74.28|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,f15a493b-b78e-43cd-a0bd-5f092c9acbc9, pattern=[METRICS]
2025-10-18T06:12:04,715 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:74.28|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:f15a493b-b78e-43cd-a0bd-5f092c9acbc9,timestamp:1760767924
2025-10-18T06:12:04,715 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79
2025-10-18T06:12:04,716 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,717 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77614.716|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,717 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:102.543|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,717 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,718 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76
2025-10-18T06:12:04,718 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,727 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,729 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924728
2025-10-18T06:12:04,730 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924
2025-10-18T06:12:04,804 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:73.63|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,e1fc061d-39b8-4681-a4d1-788df1185f0e, pattern=[METRICS]
2025-10-18T06:12:04,804 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:73.63|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:e1fc061d-39b8-4681-a4d1-788df1185f0e,timestamp:1760767924
2025-10-18T06:12:04,805 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 77
2025-10-18T06:12:04,805 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:76537.801|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:90.667|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 75

2025-10-18T06:12:04,807 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,816 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,817 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924816

2025-10-18T06:12:04,818 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,894 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.57|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,eee9304e-bd66-4109-b97c-1be74eaea35d, pattern=[METRICS]

2025-10-18T06:12:04,894 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 78

2025-10-18T06:12:04,895 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,895 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.57|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:eee9304e-bd66-4109-b97c-1be74eaea35d,timestamp:1760767924

2025-10-18T06:12:04,895 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77984.966|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,896 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:158.339|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,897 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,897 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 77

2025-10-18T06:12:04,897 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,911 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,912 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924912

2025-10-18T06:12:04,913 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,989 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.07|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,f7c77967-5d3c-4bce-a216-293f023a65b4, pattern=[METRICS]

2025-10-18T06:12:04,989 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79

2025-10-18T06:12:04,990 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.07|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:

D:f7c77967-5d3c-4bce-a216-293f023a65b4,timestamp:1760767924
2025-10-18T06:12:04,990 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,991 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77846.245|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,991 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:318.786|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,991 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,992 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 77
2025-10-18T06:12:04,992 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:05,000 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:05,001 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767925001
2025-10-18T06:12:05,002 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767925
2025-10-18T06:12:05,082 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:79.92|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767925,52d30ea3-05ce-42b9-ad51-db67b4873047, pattern=[METRICS]
2025-10-18T06:12:05,083 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 84
2025-10-18T06:12:05,083 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767925
2025-10-18T06:12:05,084 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:79.92|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:52d30ea3-05ce-42b9-ad51-db67b4873047,timestamp:1760767925
2025-10-18T06:12:05,084 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:82648.39|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767925
2025-10-18T06:12:05,085 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:104.888|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767925
2025-10-18T06:12:05,085 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767925
2025-10-18T06:12:05,089 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 82
2025-10-18T06:12:05,089 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:7.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767925
2025-10-18T06:12:05,099 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767925
2025-10-18T06:12:05,101 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767925100
2025-10-18T06:12:05,102 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767925
2025-10-18T06:12:05,178 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.95|#ModelName:m

odel,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767925,ed29189f-c539-4708-956e-52021f5657b3, pattern=[METRICS]
2025-10-18T06:12:05,179 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:05,179 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.95|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:ed29189f-c539-4708-956e-52021f5657b3,timestamp:1760767925
2025-10-18T06:12:05,180 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,180 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79840.814|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,182 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:169.132|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,182 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,183 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 78
2025-10-18T06:12:05,183 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,190 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,191 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767925190
2025-10-18T06:12:05,192 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767925
2025-10-18T06:12:05,267 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:74.67|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767925,da5a0aa4-f1e5-4faf-820c-afb1dfbca52e, pattern=[METRICS]
2025-10-18T06:12:05,267 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 78
2025-10-18T06:12:05,268 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,267 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:74.67|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:da5a0aa4-f1e5-4faf-820c-afb1dfbca52e,timestamp:1760767925
2025-10-18T06:12:05,268 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77091.449|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,269 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:91.134|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,269 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,270 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76
2025-10-18T06:12:05,270 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:02.710:[sagemaker logs]: MaxConcurrentTransforms=1, MaxPayloadInMB=6, BatchStrategy=SINGLE_RECORD

```
['torchserve', '--start', '--model-store', '/.sagemaker/ts/models', '--ts-config', '/etc/sagemaker-ts.properties', '--log-config', '/opt/conda/lib/python3.11/site-packages/sagemaker_pytorch_serving_container/etc/log4j2.xml', '--models', 'model=/opt/ml/model']
```

Warning: TorchServe is using non-default JVM parameters: -XX:-UseContainerSupport

WARNING: sun.reflect.Reflection.getCallerClass is not supported. This will impact performance.

2025-10-18T06:11:50,317 [WARN] main org.pytorch.serve.util.ConfigManager - Your torchserve instance can access any URL to load models. When deploying to production, make sure to limit the set of allowed_urls in config.properties

2025-10-18T06:11:50,321 [INFO] main org.pytorch.serve.servingsdk.impl.PluginsManager - Initializing plugins manager...

2025-10-18T06:11:50,402 [INFO] main org.pytorch.serve.metrics.configuration.MetricConfiguration - Successfully loaded metrics configuration from /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml

2025-10-18T06:11:50,499 [INFO] main org.pytorch.serve.ModelServer - Torchserve version: 0.11.0

TS Home: /opt/conda/lib/python3.11/site-packages

Current directory: /

Temp directory: /tmp

Metrics config path: /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml

Number of GPUs: 0

Number of CPUs: 2

Max heap size: 1918 M

Python executable: /opt/conda/bin/python3.11

Config file: /etc/sagemaker-ts.properties

```
['torchserve', '--start', '--model-store', '/.sagemaker/ts/models', '--ts-config', '/etc/sagemaker-ts.properties', '--log-config', '/opt/conda/lib/python3.11/site-packages/sagemaker_pytorch_serving_container/etc/log4j2.xml', '--models', 'model=/opt/ml/model']
```

Warning: TorchServe is using non-default JVM parameters: -XX:-UseContainerSupport

WARNING: sun.reflect.Reflection.getCallerClass is not supported. This will impact performance.

2025-10-18T06:11:50,317 [WARN] main org.pytorch.serve.util.ConfigManager - Your torchserve instance can access any URL to load models. When deploying to production, make sure to limit the set of allowed_urls in config.properties

2025-10-18T06:11:50,321 [INFO] main org.pytorch.serve.servingsdk.impl.PluginsManager - Initializing plugins manager...

2025-10-18T06:11:50,402 [INFO] main org.pytorch.serve.metrics.configuration.MetricConfiguration - Successfully loaded metrics configuration from /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml

2025-10-18T06:11:50,499 [INFO] main org.pytorch.serve.ModelServer - Torchserve version: 0.11.0

TS Home: /opt/conda/lib/python3.11/site-packages

Current directory: /

Temp directory: /tmp

Metrics config path: /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml

Number of GPUs: 0

Number of CPUs: 2

Max heap size: 1918 M

Python executable: /opt/conda/bin/python3.11
Config file: /etc/sagemaker-ts.properties
Inference address: http://0.0.0.0:8080
Management address: http://0.0.0.0:8080
Metrics address: http://127.0.0.1:8082
Model Store: /.sagemaker/ts/models
Initial Models: model=/opt/ml/model
Log dir: /logs
Metrics dir: /logs
Netty threads: 0
Netty client threads: 0
Inference address: http://0.0.0.0:8080
Management address: http://0.0.0.0:8080
Metrics address: http://127.0.0.1:8082
Model Store: /.sagemaker/ts/models
Initial Models: model=/opt/ml/model
Log dir: /logs
Metrics dir: /logs
Netty threads: 0
Netty client threads: 0
Default workers per model: 2
Blacklist Regex: N/A
Maximum Response Size: 6553500
Maximum Request Size: 6553500
Limit Maximum Image Pixels: true
Prefer direct buffer: false
Allowed Urls: [file://.*|http(s)?://.*]
Custom python dependency for model allowed: false
Enable metrics API: true
Metrics mode: LOG
Disable system metrics: false
Workflow Store: /.sagemaker/ts/models
CPP log config: N/A
Model config: N/A
System metrics command: default
2025-10-18T06:11:50,509 [INFO] main org.pytorch.serve.servingsdk.impl.PluginsManager - Loading snapshot serializer plugin...
2025-10-18T06:11:50,534 [INFO] main org.pytorch.serve.ModelServer - Loading initial models: /opt/ml/model
2025-10-18T06:11:50,546 [INFO] main org.pytorch.serve.archive.model.ModelArchive - createTempDir /tmp/models/99953fda636149fdbca0a2785a7884fc
2025-10-18T06:11:50,546 [INFO] main org.pytorch.serve.archive.model.ModelArchive - createSymbolicDir /tmp/models/99953fda636149fdbca0a2785a7884fc/model
2025-10-18T06:11:50,548 [WARN] main org.pytorch.serve.archive.model.ModelArchive - Model archive version is not defined. Please upgrade to torch-model-archiver 0.2.0 or higher
2025-10-18T06:11:50,548 [WARN] main org.pytorch.serve.archive.model.ModelArchive - Model archive createdOn is not defined. Please upgrade to torch-model-archiver 0.2.0 or higher
2025-10-18T06:11:50,551 [INFO] main org.pytorch.serve.wlm.ModelManager - Model model loaded.
2025-10-18T06:11:50,570 [INFO] main org.pytorch.serve.ModelServer - Initialize Inference server with: EpollServerSocketChannel.
2025-10-18T06:11:50,738 [INFO] main org.pytorch.serve.ModelServer - Inference

```
API bind to: http://0.0.0.0:8080
2025-10-18T06:11:50,739 [INFO ] main org.pytorch.serve.ModelServer - Initialize
Metrics server with: EpollServerSocketChannel.
2025-10-18T06:11:50,753 [INFO ] main org.pytorch.serve.ModelServer - Metrics AP
I bind to: http://127.0.0.1:8082
Model server started.
2025-10-18T06:11:51,484 [WARN ] pool-3-thread-1 org.pytorch.serve.metrics.Metri
cCollector - worker pid is not available yet.
Default workers per model: 2
Blacklist Regex: N/A
Maximum Response Size: 6553500
Maximum Request Size: 6553500
Limit Maximum Image Pixels: true
Prefer direct buffer: false
Allowed Urls: [file://.*|http(s)?://.*]
Custom python dependency for model allowed: false
Enable metrics API: true
Metrics mode: LOG
Disable system metrics: false
Workflow Store: /.sagemaker/ts/models
CPP log config: N/A
Model config: N/A
System metrics command: default
2025-10-18T06:11:50,509 [INFO ] main org.pytorch.serve.servingsdk.impl.PluginsM
anager - Loading snapshot serializer plugin...
2025-10-18T06:11:50,534 [INFO ] main org.pytorch.serve.ModelServer - Loading in
itial models: /opt/ml/model
2025-10-18T06:11:50,546 [INFO ] main org.pytorch.serve.archive.model.ModelArchi
ve - createTempDir /tmp/models/99953fda636149fdbca0a2785a7884fc
2025-10-18T06:11:50,546 [INFO ] main org.pytorch.serve.archive.model.ModelArchi
ve - createSymbolicDir /tmp/models/99953fda636149fdbca0a2785a7884fc/model
2025-10-18T06:11:50,548 [WARN ] main org.pytorch.serve.archive.model.ModelArchi
ve - Model archive version is not defined. Please upgrade to torch-model-archiv
er 0.2.0 or higher
2025-10-18T06:11:50,548 [WARN ] main org.pytorch.serve.archive.model.ModelArchi
ve - Model archive createdOn is not defined. Please upgrade to torch-model-arch
iver 0.2.0 or higher
2025-10-18T06:11:50,551 [INFO ] main org.pytorch.serve.wlm.ModelManager - Model
model loaded.
2025-10-18T06:11:50,570 [INFO ] main org.pytorch.serve.ModelServer - Initialize
Inference server with: EpollServerSocketChannel.
2025-10-18T06:11:50,738 [INFO ] main org.pytorch.serve.ModelServer - Inference
API bind to: http://0.0.0.0:8080
2025-10-18T06:11:50,739 [INFO ] main org.pytorch.serve.ModelServer - Initialize
Metrics server with: EpollServerSocketChannel.
2025-10-18T06:11:50,753 [INFO ] main org.pytorch.serve.ModelServer - Metrics AP
I bind to: http://127.0.0.1:8082
Model server started.
2025-10-18T06:11:51,484 [WARN ] pool-3-thread-1 org.pytorch.serve.metrics.Metri
cCollector - worker pid is not available yet.
2025-10-18T06:11:51,672 [INFO ] pool-3-thread-1 TS_METRICS - CPUUtilization.Per
cent:100.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767911
2025-10-18T06:11:51,673 [INFO ] pool-3-thread-1 TS_METRICS - DiskAvailable.Giga
bytes:47.07512283325195|#Level:Host|#hostname:0da699790d1b,timestamp:1760767911
```

2025-10-18T06:11:51,674 [INFO] pool-3-thread-1 TS_METRICS - DiskUsage.Gigabytes:8.790008544921875|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,674 [INFO] pool-3-thread-1 TS_METRICS - DiskUtilization.Percent:15.7|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,675 [INFO] pool-3-thread-1 TS_METRICS - MemoryAvailable.Megabytes:6569.18359375|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,675 [INFO] pool-3-thread-1 TS_METRICS - MemoryUsed.Megabytes:803.8046875|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,675 [INFO] pool-3-thread-1 TS_METRICS - MemoryUtilization.Percent:14.3|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,672 [INFO] pool-3-thread-1 TS_METRICS - CPUUtilization.Percent:100.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,673 [INFO] pool-3-thread-1 TS_METRICS - DiskAvailable.Gigabytes:47.07512283325195|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,674 [INFO] pool-3-thread-1 TS_METRICS - DiskUsage.Gigabytes:8.790008544921875|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,674 [INFO] pool-3-thread-1 TS_METRICS - DiskUtilization.Percent:15.7|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,675 [INFO] pool-3-thread-1 TS_METRICS - MemoryAvailable.Megabytes:6569.18359375|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,675 [INFO] pool-3-thread-1 TS_METRICS - MemoryUsed.Megabytes:803.8046875|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:51,675 [INFO] pool-3-thread-1 TS_METRICS - MemoryUtilization.Percent:14.3|#Level:Host|#hostname:0da699790dlb,timestamp:1760767911
2025-10-18T06:11:53,294 [INFO] W-9000-model_1.0-stdout MODEL_LOG - s_name_part0=/tmp/.ts.sock, s_name_part1=9000, pid=36
2025-10-18T06:11:53,295 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Listening on port: /tmp/.ts.sock.9000
2025-10-18T06:11:53,306 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Successfully loaded /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml.
2025-10-18T06:11:53,294 [INFO] W-9000-model_1.0-stdout MODEL_LOG - s_name_part0=/tmp/.ts.sock, s_name_part1=9000, pid=36
2025-10-18T06:11:53,295 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Listening on port: /tmp/.ts.sock.9000
2025-10-18T06:11:53,306 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Successfully loaded /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml.
2025-10-18T06:11:53,307 [INFO] W-9000-model_1.0-stdout MODEL_LOG - [PID]36
2025-10-18T06:11:53,307 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Torch worker started.
2025-10-18T06:11:53,308 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Python runtime: 3.11.9
2025-10-18T06:11:53,314 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Connecting to: /tmp/.ts.sock.9000
2025-10-18T06:11:53,341 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Connection accepted: /tmp/.ts.sock.9000.
2025-10-18T06:11:53,364 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767913364
2025-10-18T06:11:53,545 [INFO] W-9000-model_1.0-stdout MODEL_LOG - model_name: model, batchSize: 1
2025-10-18T06:11:53,617 [INFO] W-9001-model_1.0-stdout MODEL_LOG - s_name_part0=/tmp/.ts.sock, s_name_part1=9001, pid=37
2025-10-18T06:11:53,620 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Listening on port: /tmp/.ts.sock.9001
2025-10-18T06:11:53,630 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Successfully

```
y loaded /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml.
2025-10-18T06:11:53,631 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - [PID]37
2025-10-18T06:11:53,632 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - Torch worke
r started.
2025-10-18T06:11:53,633 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - Python runt
ime: 3.11.9
2025-10-18T06:11:53,633 [INFO ] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Connecting to: /tmp/.ts.sock.9001
2025-10-18T06:11:53,641 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - Connection
accepted: /tmp/.ts.sock.9001.
2025-10-18T06:11:53,653 [INFO ] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Looping backend response at: 1760767913653
2025-10-18T06:11:53,733 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - model_name:
model, batchSize: 1
2025-10-18T06:11:53,307 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - [PID]36
2025-10-18T06:11:53,307 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - Torch worke
r started.
2025-10-18T06:11:53,308 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - Python runt
ime: 3.11.9
2025-10-18T06:11:53,314 [INFO ] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Connecting to: /tmp/.ts.sock.9000
2025-10-18T06:11:53,341 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - Connection
accepted: /tmp/.ts.sock.9000.
2025-10-18T06:11:53,364 [INFO ] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Looping backend response at: 1760767913364
2025-10-18T06:11:53,545 [INFO ] W-9000-model_1.0-stdout MODEL_LOG - model_name:
model, batchSize: 1
2025-10-18T06:11:53,617 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - s_name_part
0=/tmp/.ts.sock, s_name_part1=9001, pid=37
2025-10-18T06:11:53,620 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - Listening o
n port: /tmp/.ts.sock.9001
2025-10-18T06:11:53,630 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - Successfull
y loaded /opt/conda/lib/python3.11/site-packages/ts/configs/metrics.yaml.
2025-10-18T06:11:53,631 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - [PID]37
2025-10-18T06:11:53,632 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - Torch worke
r started.
2025-10-18T06:11:53,633 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - Python runt
ime: 3.11.9
2025-10-18T06:11:53,633 [INFO ] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Connecting to: /tmp/.ts.sock.9001
2025-10-18T06:11:53,641 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - Connection
accepted: /tmp/.ts.sock.9001.
2025-10-18T06:11:53,653 [INFO ] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Looping backend response at: 1760767913653
2025-10-18T06:11:53,733 [INFO ] W-9001-model_1.0-stdout MODEL_LOG - model_name:
model, batchSize: 1
2025-10-18T06:11:55,863 [INFO ] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Backend response time: 2209
2025-10-18T06:11:55,864 [INFO ] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Backend response time: 2498
2025-10-18T06:11:55,864 [INFO ] W-9001-model_1.0 TS_METRICS - WorkerLoadTime.Mi
lliseconds:5299.0|#WorkerName:W-9001-model_1.0,Level:Host|#hostname:0da699790d1
b,timestamp:1760767915
2025-10-18T06:11:55,865 [INFO ] W-9000-model_1.0 TS_METRICS - WorkerLoadTime.Mi
```

lliseconds:5305.0|#WorkerName:W-9000-model_1.0,Level:Host|#hostname:0da699790d1b,timestamp:1760767915

2025-10-18T06:11:55,866 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:17.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767915

2025-10-18T06:11:55,866 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:15.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767915

2025-10-18T06:11:55,863 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 2209

2025-10-18T06:11:55,864 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 2498

2025-10-18T06:11:55,864 [INFO] W-9001-model_1.0 TS_METRICS - WorkerLoadTime.Milliseconds:5299.0|#WorkerName:W-9001-model_1.0,Level:Host|#hostname:0da699790d1b,timestamp:1760767915

2025-10-18T06:11:55,865 [INFO] W-9000-model_1.0 TS_METRICS - WorkerLoadTime.Milliseconds:5305.0|#WorkerName:W-9000-model_1.0,Level:Host|#hostname:0da699790d1b,timestamp:1760767915

2025-10-18T06:11:55,866 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:17.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767915

2025-10-18T06:11:55,866 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:15.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767915

2025-10-18T06:12:02,600 [INFO] pool-2-thread-3 ACCESS_LOG - /169.254.255.130:36970 "GET /ping HTTP/1.1" 200 32

2025-10-18T06:12:02,601 [INFO] pool-2-thread-3 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767922

2025-10-18T06:12:02,692 [INFO] epollEventLoopGroup-3-2 ACCESS_LOG - /169.254.255.130:36978 "GET /execution-parameters HTTP/1.1" 404 4

2025-10-18T06:12:02,694 [INFO] epollEventLoopGroup-3-2 TS_METRICS - Requests4XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767922

2025-10-18T06:12:02,600 [INFO] pool-2-thread-3 ACCESS_LOG - /169.254.255.130:36970 "GET /ping HTTP/1.1" 200 32

2025-10-18T06:12:02,601 [INFO] pool-2-thread-3 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767922

2025-10-18T06:12:02,692 [INFO] epollEventLoopGroup-3-2 ACCESS_LOG - /169.254.255.130:36978 "GET /execution-parameters HTTP/1.1" 404 4

2025-10-18T06:12:02,694 [INFO] epollEventLoopGroup-3-2 TS_METRICS - Requests4XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767922

2025-10-18T06:12:03,208 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,216 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923216

2025-10-18T06:12:03,219 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:03,368 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:148.52|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,eb22ad0c-d5c7-416e-854c-8a047ae4898f, pattern=[METRICS]

2025-10-18T06:12:03,371 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:148.52|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:eb22ad0c-d5c7-416e-854c-8a047ae4898f,timestamp:1760767923

2025-10-18T06:12:03,373 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 169

2025-10-18T06:12:03,374 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,375 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:157701.166|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,375 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:159.077|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,378 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,378 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 156

2025-10-18T06:12:03,379 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,450 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,451 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923451

2025-10-18T06:12:03,457 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:03,565 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:112.06|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,00d978d5-99ef-4310-9fda-9399cce91dbb, pattern=[METRICS]

2025-10-18T06:12:03,566 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 117

2025-10-18T06:12:03,567 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:112.06|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:00d978d5-99ef-4310-9fda-9399cce91dbb,timestamp:1760767923

2025-10-18T06:12:03,568 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,569 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:115682.524|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,570 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:110.69|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,208 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,216 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923216

2025-10-18T06:12:03,219 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:03,368 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:148.52|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,eb22ad0c-d5c7-416e-854c-8a047ae4898f, pattern=[METRICS]

2025-10-18T06:12:03,371 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:148.52|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:eb22ad0c-d5c7-416e-854c-8a047ae4898f,timestamp:1760767923

2025-10-18T06:12:03,373 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 169

2025-10-18T06:12:03,374 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,375 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:157701.166|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,375 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:159.077|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,378 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,378 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread read - Backend response time: 156

2025-10-18T06:12:03,379 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,450 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,451 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread read - Looping backend response at: 1760767923451

2025-10-18T06:12:03,457 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:03,565 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:112.06|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767923,00d978d5-99ef-4310-9fda-9399cce91dbb, pattern=[METRICS]

2025-10-18T06:12:03,566 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 117

2025-10-18T06:12:03,567 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:112.06|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:00d978d5-99ef-4310-9fda-9399cce91dbb,timestamp:1760767923

2025-10-18T06:12:03,568 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,569 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:115682.524|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,570 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:110.69|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,571 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,571 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread read - Backend response time: 115

2025-10-18T06:12:03,572 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:7.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,586 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,588 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread read - Looping backend response at: 1760767923588

2025-10-18T06:12:03,589 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:03,666 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:76.47|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767923,535aaf8f-12c2-4b25-aaf1-4903e128377d, pattern=[METRICS]

2025-10-18T06:12:03,667 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - Predict

ionTime.ms:76.47|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:535aaf8f-12c2-4b25-aaf1-4903e128377d,timestamp:1760767923
2025-10-18T06:12:03,667 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:03,668 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,669 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79881.416|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,670 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:118.211|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,671 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,671 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 79
2025-10-18T06:12:03,671 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:5.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,679 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,680 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923680
2025-10-18T06:12:03,681 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923
2025-10-18T06:12:03,759 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:77.19|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,139c1d02-7d60-448c-9b37-15386707603b, pattern=[METRICS]
2025-10-18T06:12:03,760 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:03,760 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,760 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:77.19|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:139c1d02-7d60-448c-9b37-15386707603b,timestamp:1760767923
2025-10-18T06:12:03,761 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79681.473|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,761 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:162.982|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,571 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,571 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 115
2025-10-18T06:12:03,572 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:7.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,586 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923
2025-10-18T06:12:03,588 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923588
2025-10-18T06:12:03,589 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend rec

eived inference at: 1760767923
2025-10-18T06:12:03,666 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:76.47|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767923,535aaf8f-12c2-4b25-aaf1-4903e128377d, pattern=[METRICS]
2025-10-18T06:12:03,667 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:76.47|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:535aaf8f-12c2-4b25-aaf1-4903e128377d,timestamp:1760767923
2025-10-18T06:12:03,667 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:03,668 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,669 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79881.416|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,670 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:118.211|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,671 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,671 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 79
2025-10-18T06:12:03,671 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:5.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,679 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,680 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923680
2025-10-18T06:12:03,681 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923
2025-10-18T06:12:03,759 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:77.19|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767923,139c1d02-7d60-448c-9b37-15386707603b, pattern=[METRICS]
2025-10-18T06:12:03,760 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:03,760 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,760 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:77.19|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:139c1d02-7d60-448c-9b37-15386707603b,timestamp:1760767923
2025-10-18T06:12:03,761 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79681.473|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,761 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:162.982|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,762 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,762 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 78
2025-10-18T06:12:03,762 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,775 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,776 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923776

2025-10-18T06:12:03,778 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:03,855 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:76.6|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767923,55b975c1-dc4b-46b1-92fc-af89488e5ce7, pattern=[METRICS]

2025-10-18T06:12:03,856 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81

2025-10-18T06:12:03,857 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,857 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:76.6|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:55b975c1-dc4b-46b1-92fc-af89488e5ce7,timestamp:1760767923

2025-10-18T06:12:03,858 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79770.92|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,858 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:104.25|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,859 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,860 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 79

2025-10-18T06:12:03,862 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:7.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,870 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,870 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923870

2025-10-18T06:12:03,873 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:03,762 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,762 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 78

2025-10-18T06:12:03,762 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,775 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767923

2025-10-18T06:12:03,776 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923776

2025-10-18T06:12:03,778 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:03,855 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:76.6|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767923,55b975c1-dc4b-46b1-92fc-af89488e5ce7, pattern=[METRICS]

2025-10-18T06:12:03,856 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.13
0:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:03,857 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Coun
t:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,857 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - Predict
ionTime.ms:76.6|#modelName:model,Level:Model|#hostname:0da699790dlb,requestID:5
5b975c1-dc4b-46b1-92fc-af89488e5ce7,timestamp:1760767923
2025-10-18T06:12:03,858 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_late
ncy_microseconds.Microseconds:79770.92|#model_name:model,model_version:default
t|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,858 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latenc
y_microseconds.Microseconds:104.25|#model_name:model,model_version:default|#hos
tname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,859 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Millise
conds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,860 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Backend response time: 79
2025-10-18T06:12:03,862 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTim
e.Milliseconds:7.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,870 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inferen
ce_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0
da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,870 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Looping backend response at: 1760767923870
2025-10-18T06:12:03,873 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend rec
eived inference at: 1760767923
2025-10-18T06:12:03,955 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.W
orkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:82.16|#modelName:m
odel,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767923,4416b3f4-bd1b-44
85-bc1d-448e7bf30711, pattern=[METRICS]
2025-10-18T06:12:03,955 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.13
0:36984 "POST /invocations HTTP/1.1" 200 86
2025-10-18T06:12:03,956 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Coun
t:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,956 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - Predict
ionTime.ms:82.16|#modelName:model,Level:Model|#hostname:0da699790dlb,requestI
D:4416b3f4-bd1b-4485-bc1d-448e7bf30711,timestamp:1760767923
2025-10-18T06:12:03,956 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_late
ncy_microseconds.Microseconds:85149.584|#model_name:model,model_version:default
t|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,957 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latenc
y_microseconds.Microseconds:107.279|#model_name:model,model_version:default|#ho
stname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,957 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Millise
conds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,957 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Backend response time: 84
2025-10-18T06:12:03,958 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTim
e.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,967 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inferen
ce_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0
da699790dlb,timestamp:1760767923
2025-10-18T06:12:03,968 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Looping backend response at: 1760767923968

2025-10-18T06:12:03,974 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:04,052 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:78.89|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,be1244e1-4fb2-4c32-8c4c-2e9803e5e3e9, pattern=[METRICS]

2025-10-18T06:12:04,053 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 87

2025-10-18T06:12:04,053 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,054 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:85517.614|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,054 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:97.436|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,055 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,055 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 84

2025-10-18T06:12:04,055 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,054 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:78.89|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:be1244e1-4fb2-4c32-8c4c-2e9803e5e3e9,timestamp:1760767924

2025-10-18T06:12:04,069 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:03,955 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:82.16|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767923,4416b3f4-bd1b-4485-bc1d-448e7bf30711, pattern=[METRICS]

2025-10-18T06:12:03,955 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 86

2025-10-18T06:12:03,956 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,956 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:82.16|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:4416b3f4-bd1b-4485-bc1d-448e7bf30711,timestamp:1760767923

2025-10-18T06:12:03,956 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:85149.584|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,957 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:107.279|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,957 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,957 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 84

2025-10-18T06:12:03,958 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,967 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767923

2025-10-18T06:12:03,968 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767923968

2025-10-18T06:12:03,974 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767923

2025-10-18T06:12:04,052 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:78.89|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,be1244e1-4fb2-4c32-8c4c-2e9803e5e3e9, pattern=[METRICS]

2025-10-18T06:12:04,053 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 87

2025-10-18T06:12:04,053 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,054 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:85517.614|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,054 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:97.436|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,055 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,055 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 84

2025-10-18T06:12:04,055 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,054 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:78.89|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:be1244e1-4fb2-4c32-8c4c-2e9803e5e3e9,timestamp:1760767924

2025-10-18T06:12:04,069 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,070 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924070

2025-10-18T06:12:04,073 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.59|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,6463fc01-0332-4cde-a248-a3d8a9620b7a, pattern=[METRICS]

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.59|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:6463fc01-0332-4cde-a248-a3d8a9620b7a,timestamp:1760767924

2025-10-18T06:12:04,150 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77731.709|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,150 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:103.801|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,150 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,151 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread

read - Backend response time: 78
2025-10-18T06:12:04,151 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:3.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,165 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,167 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924167
2025-10-18T06:12:04,169 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924
2025-10-18T06:12:04,245 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:76.2|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,db313afb-9d9e-4212-a11c-6b71df7d258b, pattern=[METRICS]
2025-10-18T06:12:04,245 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 80
2025-10-18T06:12:04,246 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,246 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:76.2|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:db313afb-9d9e-4212-a11c-6b71df7d258b,timestamp:1760767924
2025-10-18T06:12:04,247 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79525.992|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,247 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:220.372|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,248 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,248 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76
2025-10-18T06:12:04,248 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,260 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,261 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924261
2025-10-18T06:12:04,262 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924
2025-10-18T06:12:04,339 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.89|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,b5b61720-5375-4580-b0c8-2b0354434701, pattern=[METRICS]
2025-10-18T06:12:04,339 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79
2025-10-18T06:12:04,340 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,070 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924070
2025-10-18T06:12:04,073 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924
2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.59|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,6463fc01-0332-4cde-a248-a3d8a9620b7a, pattern=[METRICS]

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,148 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.59|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:6463fc01-0332-4cde-a248-a3d8a9620b7a,timestamp:1760767924

2025-10-18T06:12:04,150 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77731.709|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,150 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:103.801|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,150 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,151 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 78

2025-10-18T06:12:04,151 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:3.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,165 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,167 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924167

2025-10-18T06:12:04,169 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,245 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:76.2|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,db313afb-9d9e-4212-a11c-6b71df7d258b, pattern=[METRICS]

2025-10-18T06:12:04,245 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 80

2025-10-18T06:12:04,246 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,246 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:76.2|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:db313afb-9d9e-4212-a11c-6b71df7d258b,timestamp:1760767924

2025-10-18T06:12:04,247 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79525.992|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,247 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:220.372|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,248 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,248 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76

2025-10-18T06:12:04,248 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,260 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,261 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924261

2025-10-18T06:12:04,262 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,339 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.89|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,b5b61720-5375-4580-b0c8-2b0354434701, pattern=[METRICS]

2025-10-18T06:12:04,339 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79

2025-10-18T06:12:04,340 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,340 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.89|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:b5b61720-5375-4580-b0c8-2b0354434701,timestamp:1760767924

2025-10-18T06:12:04,341 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77997.837|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,342 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:221.097|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,342 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,342 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 77

2025-10-18T06:12:04,343 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:5.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,353 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,354 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924354

2025-10-18T06:12:04,355 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,437 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:81.17|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,876916ed-83a6-47bc-99eb-c38dcedf4ca2, pattern=[METRICS]

2025-10-18T06:12:04,437 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:81.17|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:876916ed-83a6-47bc-99eb-c38dcedf4ca2,timestamp:1760767924

2025-10-18T06:12:04,440 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 85

2025-10-18T06:12:04,440 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:84146.256|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:97.85|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread

read - Backend response time: 84
2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:3.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,454 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,456 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924456
2025-10-18T06:12:04,461 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924
2025-10-18T06:12:04,534 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.86|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,2c08a9c5-b823-4cf7-b907-546f89204feb, pattern=[METRICS]
2025-10-18T06:12:04,534 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.86|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:2c08a9c5-b823-4cf7-b907-546f89204feb,timestamp:1760767924
2025-10-18T06:12:04,535 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:04,536 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,536 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_lateness_microseconds.Microseconds:79458.604|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,536 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:113.16|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,537 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,537 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 78
2025-10-18T06:12:04,537 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,340 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.89|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:b5b61720-5375-4580-b0c8-2b0354434701,timestamp:1760767924
2025-10-18T06:12:04,341 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_lateness_microseconds.Microseconds:77997.837|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,342 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:221.097|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,342 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,342 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 77
2025-10-18T06:12:04,343 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:5.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,353 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,354 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924354
2025-10-18T06:12:04,355 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend rec

eived inference at: 1760767924
2025-10-18T06:12:04,437 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:81.17|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,876916ed-83a6-47bc-99eb-c38dcedf4ca2, pattern=[METRICS]
2025-10-18T06:12:04,437 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:81.17|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:876916ed-83a6-47bc-99eb-c38dcedf4ca2,timestamp:1760767924
2025-10-18T06:12:04,440 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 85
2025-10-18T06:12:04,440 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:84146.256|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:97.85|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 84
2025-10-18T06:12:04,441 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:3.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,454 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,456 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924456
2025-10-18T06:12:04,461 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924
2025-10-18T06:12:04,534 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.86|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,2c08a9c5-b823-4cf7-b907-546f89204feb, pattern=[METRICS]
2025-10-18T06:12:04,534 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.86|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:2c08a9c5-b823-4cf7-b907-546f89204feb,timestamp:1760767924
2025-10-18T06:12:04,535 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:04,536 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,536 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79458.604|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,536 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:113.16|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,537 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924
2025-10-18T06:12:04,537 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 78
2025-10-18T06:12:04,537 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,546 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,548 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924548

2025-10-18T06:12:04,550 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,625 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.4|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,99cf3977-ef89-4533-91d6-d8c24f50c4f3, pattern=[METRICS]

2025-10-18T06:12:04,626 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 80

2025-10-18T06:12:04,626 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.4|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:99cf3977-ef89-4533-91d6-d8c24f50c4f3,timestamp:1760767924

2025-10-18T06:12:04,627 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,627 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79242.782|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,628 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:110.513|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,628 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,629 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76

2025-10-18T06:12:04,629 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,637 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,639 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924638

2025-10-18T06:12:04,640 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,714 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:74.28|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,f15a493b-b78e-43cd-a0bd-5f092c9acbc9, pattern=[METRICS]

2025-10-18T06:12:04,715 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:74.28|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:f15a493b-b78e-43cd-a0bd-5f092c9acbc9,timestamp:1760767924

2025-10-18T06:12:04,715 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79

2025-10-18T06:12:04,716 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,717 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77614.716|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,717 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:102.543|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,717 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,718 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76

2025-10-18T06:12:04,718 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,727 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,729 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924728

2025-10-18T06:12:04,730 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,804 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:73.63|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,e1fc061d-39b8-4681-a4d1-788df1185f0e, pattern=[METRICS]

2025-10-18T06:12:04,804 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:73.63|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:e1fc061d-39b8-4681-a4d1-788df1185f0e,timestamp:1760767924

2025-10-18T06:12:04,805 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 77

2025-10-18T06:12:04,805 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,546 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,548 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924548

2025-10-18T06:12:04,550 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,625 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.4|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,99cf3977-ef89-4533-91d6-d8c24f50c4f3, pattern=[METRICS]

2025-10-18T06:12:04,626 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 80

2025-10-18T06:12:04,626 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.4|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:99cf3977-ef89-4533-91d6-d8c24f50c4f3,timestamp:1760767924

2025-10-18T06:12:04,627 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,627 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79242.782|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,628 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:110.513|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,628 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,629 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76

2025-10-18T06:12:04,629 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,637 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,639 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924638

2025-10-18T06:12:04,640 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,714 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:74.28|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,f15a493b-b78e-43cd-a0bd-5f092c9acbc9, pattern=[METRICS]

2025-10-18T06:12:04,715 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:74.28|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:f15a493b-b78e-43cd-a0bd-5f092c9acbc9,timestamp:1760767924

2025-10-18T06:12:04,715 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79

2025-10-18T06:12:04,716 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,717 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77614.716|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,717 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:102.543|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,717 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,718 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76

2025-10-18T06:12:04,718 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,727 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,729 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924728

2025-10-18T06:12:04,730 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,804 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:73.63|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,e1fc061d-39b8-4681-a4d1-788df1185f0e, pattern=[METRICS]

2025-10-18T06:12:04,804 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:73.63|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:e1fc061d-39b8-4681-a4d1-788df1185f0e,timestamp:1760767924

2025-10-18T06:12:04,805 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 77

2025-10-18T06:12:04,805 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:76537.801|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:90.667|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 75

2025-10-18T06:12:04,807 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,816 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,817 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924816

2025-10-18T06:12:04,818 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,894 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.57|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767924,eee9304e-bd66-4109-b97c-1be74eaea35d, pattern=[METRICS]

2025-10-18T06:12:04,894 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 78

2025-10-18T06:12:04,895 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,895 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.57|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:eee9304e-bd66-4109-b97c-1be74eaea35d,timestamp:1760767924

2025-10-18T06:12:04,895 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77984.966|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,896 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:158.339|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,897 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,897 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 77

2025-10-18T06:12:04,897 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,911 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,912 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767924912

2025-10-18T06:12:04,913 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:76537.801|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:90.667|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,806 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 75

2025-10-18T06:12:04,807 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTim

e.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,816 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inferen
ce_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0
da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,817 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Looping backend response at: 1760767924816
2025-10-18T06:12:04,818 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend rec
eived inference at: 1760767924
2025-10-18T06:12:04,894 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.W
orkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.57|#ModelName:m
odel,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,eee9304e-bd66-41
09-b97c-1be74eaea35d, pattern=[METRICS]
2025-10-18T06:12:04,894 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.13
0:36984 "POST /invocations HTTP/1.1" 200 78
2025-10-18T06:12:04,895 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Coun
t:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,895 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - Predict
ionTime.ms:75.57|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestI
D:eee9304e-bd66-4109-b97c-1be74eaea35d,timestamp:1760767924
2025-10-18T06:12:04,895 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_late
ncy_microseconds.Microseconds:77984.966|#model_name:model,model_version:default
|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,896 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latenc
y_microseconds.Microseconds:158.339|#model_name:model,model_version:default|#ho
stname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,897 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Millise
conds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,897 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Backend response time: 77
2025-10-18T06:12:04,897 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTim
e.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,911 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inferen
ce_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0
da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,912 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerTh
read - Looping backend response at: 1760767924912
2025-10-18T06:12:04,913 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend rec
eived inference at: 1760767924
2025-10-18T06:12:04,989 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.W
orkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.07|#ModelName:m
odel,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,f7c77967-5d3c-4b
ce-a216-293f023a65b4, pattern=[METRICS]
2025-10-18T06:12:04,989 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.W
orkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.07|#ModelName:m
odel,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767924,f7c77967-5d3c-4b
ce-a216-293f023a65b4, pattern=[METRICS]
2025-10-18T06:12:04,989 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.13
0:36984 "POST /invocations HTTP/1.1" 200 79
2025-10-18T06:12:04,990 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - Predict
ionTime.ms:75.07|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestI
D:f7c77967-5d3c-4bce-a216-293f023a65b4,timestamp:1760767924
2025-10-18T06:12:04,990 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Coun
t:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924
2025-10-18T06:12:04,991 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_late

ncy_microseconds.Microseconds:77846.245|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,991 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:318.786|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,991 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:04,992 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 77

2025-10-18T06:12:04,992 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:05,000 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767924

2025-10-18T06:12:05,001 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767925001

2025-10-18T06:12:05,002 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767925

2025-10-18T06:12:05,082 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:79.92|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767925,52d30ea3-05ce-42b9-ad51-db67b4873047, pattern=[METRICS]

2025-10-18T06:12:05,083 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 84

2025-10-18T06:12:05,083 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767925

2025-10-18T06:12:05,084 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:79.92|#ModelName:model,Level:Model|#hostname:0da699790dlb,requestID:52d30ea3-05ce-42b9-ad51-db67b4873047,timestamp:1760767925

2025-10-18T06:12:05,084 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:82648.39|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767925

2025-10-18T06:12:05,085 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:104.888|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767925

2025-10-18T06:12:05,085 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767925

2025-10-18T06:12:05,089 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 82

2025-10-18T06:12:05,089 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:7.0|#Level:Host|#hostname:0da699790dlb,timestamp:1760767925

2025-10-18T06:12:05,099 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790dlb,timestamp:1760767925

2025-10-18T06:12:05,101 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767925100

2025-10-18T06:12:05,102 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767925

2025-10-18T06:12:05,178 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.95|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790dlb,1760767925,ed29189f-c539-4708-956e-52021f5657b3, pattern=[METRICS]

2025-10-18T06:12:05,179 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 81

2025-10-18T06:12:05,179 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.95|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:ed29189f-c539-4708-956e-52021f5657b3,timestamp:1760767925

2025-10-18T06:12:05,180 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,180 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79840.814|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,182 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:169.132|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,182 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,183 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 78

2025-10-18T06:12:05,183 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,190 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,191 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767925190

2025-10-18T06:12:05,192 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767925

2025-10-18T06:12:05,267 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:74.67|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767925,da5a0aa4-f1e5-4faf-820c-afb1dfbca52e, pattern=[METRICS]

2025-10-18T06:12:05,267 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 78

2025-10-18T06:12:05,268 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,267 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:74.67|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:da5a0aa4-f1e5-4faf-820c-afb1dfbca52e,timestamp:1760767925

2025-10-18T06:12:05,268 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77091.449|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,269 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:91.134|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,269 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,270 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76

2025-10-18T06:12:05,270 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:04,989 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 79

2025-10-18T06:12:04,990 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.07|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:f7c77967-5d3c-4bce-a216-293f023a65b4,timestamp:1760767924

2025-10-18T06:12:04,990 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,991 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77846.245|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,991 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:318.786|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,991 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:04,992 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 77

2025-10-18T06:12:04,992 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:05,000 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767924

2025-10-18T06:12:05,001 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767925001

2025-10-18T06:12:05,002 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767925

2025-10-18T06:12:05,082 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:79.92|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767925,52d30ea3-05ce-42b9-ad51-db67b4873047, pattern=[METRICS]

2025-10-18T06:12:05,083 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.130:36984 "POST /invocations HTTP/1.1" 200 84

2025-10-18T06:12:05,083 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,084 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:79.92|#ModelName:model,Level:Model|#hostname:0da699790d1b,requestID:52d30ea3-05ce-42b9-ad51-db67b4873047,timestamp:1760767925

2025-10-18T06:12:05,084 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:82648.39|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,085 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:104.888|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,085 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,089 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 82

2025-10-18T06:12:05,089 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:7.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,099 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925

2025-10-18T06:12:05,101 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767925100

2025-10-18T06:12:05,102 [INFO] W-9000-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767925

2025-10-18T06:12:05,178 [INFO] W-9000-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:75.95|#ModelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767925,ed29189f-c539-4708-956e-52021f5657b3, pattern=[METRICS]

2025-10-18T06:12:05,179 [INFO] W-9000-model_1.0 ACCESS_LOG - /169.254.255.13

0:36984 "POST /invocations HTTP/1.1" 200 81
2025-10-18T06:12:05,179 [INFO] W-9000-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:75.95|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:ed29189f-c539-4708-956e-52021f5657b3,timestamp:1760767925
2025-10-18T06:12:05,180 [INFO] W-9000-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,180 [INFO] W-9000-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:79840.814|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,182 [INFO] W-9000-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:169.132|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,182 [INFO] W-9000-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,183 [INFO] W-9000-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 78
2025-10-18T06:12:05,183 [INFO] W-9000-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:6.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,190 [INFO] epollEventLoopGroup-3-3 TS_METRICS - ts_inference_requests_total.Count:1.0|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,191 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Looping backend response at: 1760767925190
2025-10-18T06:12:05,192 [INFO] W-9001-model_1.0-stdout MODEL_LOG - Backend received inference at: 1760767925
2025-10-18T06:12:05,267 [INFO] W-9001-model_1.0-stdout org.pytorch.serve.wlm.WorkerLifeCycle - result=[METRICS]PredictionTime.Milliseconds:74.67|#modelName:model,Level:Model|#type:GAUGE|#hostname:0da699790d1b,1760767925,da5a0aa4-f1e5-4faf-820c-afb1dfbca52e, pattern=[METRICS]
2025-10-18T06:12:05,267 [INFO] W-9001-model_1.0 ACCESS_LOG - /169.254.255.13
0:36984 "POST /invocations HTTP/1.1" 200 78
2025-10-18T06:12:05,268 [INFO] W-9001-model_1.0 TS_METRICS - Requests2XX.Count:1.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,267 [INFO] W-9001-model_1.0-stdout MODEL_METRICS - PredictionTime.ms:74.67|#modelName:model,Level:Model|#hostname:0da699790d1b,requestID:da5a0aa4-f1e5-4faf-820c-afb1dfbca52e,timestamp:1760767925
2025-10-18T06:12:05,268 [INFO] W-9001-model_1.0 TS_METRICS - ts_inference_latency_microseconds.Microseconds:77091.449|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,269 [INFO] W-9001-model_1.0 TS_METRICS - ts_queue_latency_microseconds.Microseconds:91.134|#model_name:model,model_version:default|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,269 [INFO] W-9001-model_1.0 TS_METRICS - QueueTime.Milliseconds:0.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:05,270 [INFO] W-9001-model_1.0 org.pytorch.serve.wlm.WorkerThread - Backend response time: 76
2025-10-18T06:12:05,270 [INFO] W-9001-model_1.0 TS_METRICS - WorkerThreadTime.Milliseconds:4.0|#Level:Host|#hostname:0da699790d1b,timestamp:1760767925
2025-10-18T06:12:02.710:[sagemaker logs]: MaxConcurrentTransforms=1, MaxPayloadInMB=6, BatchStrategy=SINGLE_RECORD
Batch input : s3://sagemaker-us-east-1-533266958221/batch/input/batch_input.jsonl
Batch output: s3://sagemaker-us-east-1-533266958221/batch/output/20251018-060631/

In []: